

```
In [1]: #Overall IPC Crime Category
```

```
In [2]: import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from linearmodels.panel import PanelOLS
from linearmodels.panel import compare
```

```
In [3]: df = pd.read_csv('panel_data/new data.csv')
df['l_ipc'] = np.log(df['ipc_cr'])
df.head()
```

```
Out[3]:   s.no. districts year ipc_crimes pop_in_lak ipc_cr avg_temp tot_rf l_ipc
0      1    ariyalur 2011      1953     7.52  259.54  28.312353 1103.207404 5.558911
1      1    ariyalur 2012      2480     7.63  324.83  28.777312  973.207972 5.783302
2      1    ariyalur 2013      1991     7.76  256.53  28.730311  870.158045 5.547246
3      1    ariyalur 2014      1615     7.88  204.88  28.536042 1090.802339 5.322424
4      1    ariyalur 2015      1386     8.00  173.00  28.565911 1501.644532 5.153292
```

```
In [4]: df = df.set_index(['districts','year'])
y = df['l_ipc']
X = df[['avg_temp','tot_rf']]
```

```
In [5]: #PooledOLS Estimation
X = sm.add_constant(X)
pols = PanelOLS(y,X)
pols_result = pols.fit()
print(pols_result.summary)
```

PanelOLS Estimation Summary

Dep. Variable:	l_ipc	R-squared:	0.0457
Estimator:	PanelOLS	R-squared (Between):	0.0408
No. Observations:	384	R-squared (Within):	0.0466
Date:	Tue, Nov 11 2025	R-squared (Overall):	0.0457
Time:	15:29:21	Log-likelihood	-265.57
Cov. Estimator:	Unadjusted	F-statistic:	9.1189
Entities:	32	P-value	0.0001
Avg Obs:	12.000	Distribution:	F(2,381)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	9.1189
		P-value	0.0001
Time periods:	12	Distribution:	F(2,381)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	4.8315	0.2463	19.618	0.0000	4.3473	5.3158
avg_temp	0.0187	0.0078	2.3838	0.0176	0.0033	0.0341
tot_rf	0.0002	5.757e-05	4.2059	0.0000	0.0001	0.0004

```
In [6]: #FE Model Estimation
X = sm.add_constant(X)
FEmodel = PanelOLS(y,X,entity_effects=True)
feresult = FEmodel.fit()
print(feresult.summary)
```

PanelOLS Estimation Summary

Dep. Variable:	l_ipc	R-squared:	0.0474
Estimator:	PanelOLS	R-squared (Between):	0.0275
No. Observations:	384	R-squared (Within):	0.0474
Date:	Tue, Nov 11 2025	R-squared (Overall):	0.0444
Time:	15:29:21	Log-likelihood	-233.43
Cov. Estimator:	Unadjusted	F-statistic:	8.7101
Entities:	32	P-value	0.0002
Avg Obs:	12.000	Distribution:	F(2, 350)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	8.7101
		P-value	0.0002
Time periods:	12	Distribution:	F(2, 350)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	4.6751	1.1259	4.1523	0.0000	2.4607	6.8895
avg_temp	0.0229	0.0403	0.5669	0.5712	-0.0564	0.1022
tot_rf	0.0003	7.078e-05	3.9602	0.0001	0.0001	0.0004

F-test for Poolability: 2.0570

P-value: 0.0010

Distribution: F(31,350)

Included effects: Entity

```
In [9]: #RE Model Estimation
from linearmodels.panel import RandomEffects
import statsmodels.api as sm
X = sm.add_constant(X)
REmodel = RandomEffects(y,X)
reresult = REmodel.fit()
print(reresult.summary)
```

RandomEffects Estimation Summary

Dep. Variable:	l_ipc	R-squared:	0.0463
Estimator:	RandomEffects	R-squared (Between):	0.0362
No. Observations:	384	R-squared (Within):	0.0471
Date:	Tue, Nov 11 2025	R-squared (Overall):	0.0455
Time:	15:29:22	Log-likelihood	-249.01
Cov. Estimator:	Unadjusted	F-statistic:	9.2495
Entities:	32	P-value	0.0001
Avg Obs:	12.000	Distribution:	F(2,381)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	9.2495
		P-value	0.0001
Time periods:	12	Distribution:	F(2,381)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	4.7907	0.3184	15.046	0.0000	4.1646	5.4168
avg_temp	0.0195	0.0106	1.8415	0.0663	-0.0013	0.0403
tot_rf	0.0003	6.03e-05	4.2891	0.0000	0.0001	0.0004

In [10]:

```
#Hausman Test
from numpy.linalg import inv
from scipy.stats import chi2

b_FE = feresult.params
b_RE = reresult.params

common_coef = list(set(b_FE.index) & set(b_RE.index))

if 'const' in common_coef:
    common_coef.remove('const')

b_FE = b_FE[common_coef]
b_RE = b_RE[common_coef]

V_FE = feresult.cov
V_RE = reresult.cov

diff = b_FE - b_RE
diff_var = V_FE.loc[common_coef, common_coef] - V_RE.loc[common_coef, common_coef]

hausman_stat = np.dot(np.dot(diff.T, inv(diff_var)), diff)

df_h = len(diff)
p_value = 1 - chi2.cdf(hausman_stat, df_h)

print("Hausman Test Statistic:", round(hausman_stat, 3))
print("Degrees of Freedom:", df_h)
```

```
print("p-value:", round(p_value, 4))
```

```
Hausman Test Statistic: 0.567  
Degrees of Freedom: 2  
p-value: 0.753
```

In [11]: *#Diagnostic Checks*

```
from statsmodels.stats.diagnostic import het_breushpagan, het_white  
from statsmodels.stats.stattools import durbin_watson
```

In [12]: *#Test for Heteroskedasticity*

```
#H0: No heteroskedasticity  
#H1: Heteroskedasticity exists
```

```
#p-value <= 0.05 ---> Heteroskedasticity; p-value > 0.05 ---> Homoskedasticity
```

```
print('Breusch-Pagan Test')  
residuals = reresult.resids  
bp_test = het_breushpagan(residuals, X)  
bp_labels = ['Lagrange multiplier statistic', 'p-value', 'f-value', 'f p-value']  
print(dict(zip(bp_labels, bp_test)))  
print()  
print('White Test')  
white_test = het_white(residuals, X)  
white_labels = ['LM stat', 'LM p-value', 'F p-value']  
print(dict(zip(white_labels, white_test)))
```

Breusch-Pagan Test

```
{'Lagrange multiplier statistic': np.float64(9.98334328950044), 'p-value': np.float64(0.006794297342475483), 'f-value': np.float64(5.084872190924652), 'f p-value': np.float64(0.006616473096637813)}
```

White Test

```
{'LM stat': np.float64(13.545411746797555), 'LM p-value': np.float64(0.018770041534266983), 'F p-value': np.float64(2.7642608852180715)}
```

In [13]: *#Test for serial correlation (autocorrelation)*

```
#Durbin-Watson statistic ranges between 0 to 4
```

```
#DW statistic = 2 ---> No autocorrelation  
#DW statistic < 2 ---> Positive autocorrelation  
#DW statistic > 2 ---> Negative autocorrelation
```

```
print('Durbin-Watson Test')  
dw_value = durbin_watson(residuals)  
print("Durbin-Watson statistic: ", round(dw_value, 3))
```

Durbin-Watson Test

```
Durbin-Watson statistic: 1.701
```

In [14]: *from scipy import stats*

```
#Test for cross-section dependency
```

```
#H0: No cross-section dependency
```

```
#H1: Cross-section dependency exists

print('Breusch-Pagan LM Test')
resid_df = residuals.unstack(level=0)
T = resid_df.shape[0]
N = resid_df.shape[1]

rho = resid_df.corr().values
upper_tri_idx = np.triu_indices(N, k=1)
rho_upper = rho[upper_tri_idx]
LM_stat = T * np.sum(rho_upper**2)
p_value = 1 - stats.chi2.cdf(LM_stat, N*(N-1)/2)

print(f"Breusch-Pagan LM statistic: {LM_stat:.3f}")
print(f"p-value: {p_value:.4f}")
print()

print('Pearson CD Test')
CD_stat = np.sqrt(2 / (N*(N-1))) * np.sum(rho_upper)
p_value_cd = 2 * (1 - stats.norm.cdf(abs(CD_stat)))

print(f"Pesaran CD statistic: {CD_stat:.3f}")
print(f"p-value: {p_value_cd:.4f}")
```

```
Breusch-Pagan LM Test
Breusch-Pagan LM statistic: 3117.798
p-value: 0.0000
```

```
Pearson CD Test
Pesaran CD statistic: 14.825
p-value: 0.0000
```

```
In [15]: #Re-estimate RE Model
```

```
In [16]: #RE with cov.type 'clustered'
re_model_robust1 = REmodel.fit(cov_type='clustered', cluster_entity=True)
print(re_model_robust1.summary)
```

RandomEffects Estimation Summary

Dep. Variable:	l_ipc	R-squared:	0.0463
Estimator:	RandomEffects	R-squared (Between):	0.0362
No. Observations:	384	R-squared (Within):	0.0471
Date:	Tue, Nov 11 2025	R-squared (Overall):	0.0455
Time:	15:29:22	Log-likelihood	-249.01
Cov. Estimator:	Clustered	F-statistic:	9.2495
Entities:	32	P-value	0.0001
Avg Obs:	12.000	Distribution:	F(2,381)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	15.015
		P-value	0.0000
Time periods:	12	Distribution:	F(2,381)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	4.7907	0.4626	10.357	0.0000	3.8812	5.7002
avg_temp	0.0195	0.0160	1.2148	0.2252	-0.0121	0.0510
tot_rf	0.0003	4.723e-05	5.4757	0.0000	0.0002	0.0004

```
In [17]: #RE with cov.type 'kernel' (Driscoll-Kraay Method)
re_model_robust2 = REmodel.fit(cov_type='kernel')
print(re_model_robust2.summary)
```

RandomEffects Estimation Summary						
Dep. Variable:	l_ipc	R-squared:	0.0463			
Estimator:	RandomEffects	R-squared (Between):	0.0362			
No. Observations:	384	R-squared (Within):	0.0471			
Date:	Tue, Nov 11 2025	R-squared (Overall):	0.0455			
Time:	15:29:22	Log-likelihood	-249.01			
Cov. Estimator:	Driscoll-Kraay					
		F-statistic:	9.2495			
Entities:	32	P-value	0.0001			
Avg Obs:	12.000	Distribution:	F(2,381)			
Min Obs:	12.000					
Max Obs:	12.000	F-statistic (robust):	4.6172			
		P-value	0.0104			
Time periods:	12	Distribution:	F(2,381)			
Avg Obs:	32.000					
Min Obs:	32.000					
Max Obs:	32.000					
Parameter Estimates						
	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	4.7907	0.2095	22.865	0.0000	4.3787	5.2027
avg_temp	0.0195	0.0064	3.0323	0.0026	0.0068	0.0321
tot_rf	0.0003	0.0001	2.1170	0.0349	1.842e-05	0.0005

```
In [18]: # Check residuals and fitted values
df['residuals1'] = re_model_robust1.resids
df['fitted1'] = re_model_robust1.fitted_values

import matplotlib.pyplot as plt

plt.scatter(df['fitted1'], df['residuals1'], alpha=0.6)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values (RE model with Clustered Standard Errors)')
plt.show()

sm.qqplot(df['residuals1'], line='45', fit=True)
plt.title('Q-Q Plot of Residuals')
plt.show()

plt.hist(df['residuals1'], bins=30, edgecolor='black', alpha=0.7)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (RE model)')
plt.show()


resid_df = df['residuals1'].unstack(level=0)
plt.plot(resid_df.mean(axis=1))
plt.title('Average Residuals over Time')
plt.xlabel('Year')
plt.ylabel('Mean Residual')
```



```
plt.show()


from scipy.stats import shapiro

#Test for normality

stat, p = shapiro(df['residuals1'])
print(f"Shapiro-Wilk Test: Statistic={stat:.3f}, p-value={p:.4f}")
```







Shapiro-Wilk Test: Statistic=0.797, p-value=0.0000

```
In [19]: # Check residuals and fitted values
df['residuals2'] = re_model_robust2.resids
df['fitted2'] = re_model_robust2.fitted_values

import matplotlib.pyplot as plt

plt.scatter(df['fitted2'], df['residuals2'], alpha=0.6)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values (RE model with Driscoll-Kraay)')
plt.show()

sm.qqplot(df['residuals2'], line='45', fit=True)
plt.title('Q-Q Plot of Residuals')
plt.show()

plt.hist(df['residuals2'], bins=30, edgecolor='black', alpha=0.7)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (RE model)')
plt.show()


resid_df = df['residuals2'].unstack(level=0)
plt.plot(resid_df.mean(axis=1))
plt.title('Average Residuals over Time')
plt.xlabel('Year')
```



```
plt.ylabel('Mean Residual')
plt.show()


from scipy.stats import shapiro

#Test for normality

stat, p = shapiro(df['residuals2'])
print(f"Shapiro-Wilk Test: Statistic={stat:.3f}, p-value={p:.4f}")
```







Shapiro-Wilk Test: Statistic=0.797, p-value=0.0000

In []: