

```
In [1]: #Violent Crime Category
```

```
In [2]: import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from linearmodels.panel import PanelOLS
from linearmodels.panel import compare
```

```
In [3]: df = pd.read_csv('panel_data/violent_new.csv')
#df['L_ipc'] = np.log(df['ipc_cr'])
df.head()
```

```
Out[3]:   s.no. districts year type violent_crimes pop_in_lak violent_cr avg_temp tot_rf
0      1    ariyalur 2011 violent crimes        148     7.52    19.67  28.312353 1103.207
1      1    ariyalur 2012 violent crimes        176     7.63    23.05  28.777312  973.207
2      1    ariyalur 2013 violent crime       155     7.76    19.97  28.730311  870.158
3      1    ariyalur 2014 violent crime       127     7.88    16.11  28.536042 1090.802
4      1    ariyalur 2015 violent crime        85     8.00    10.60  28.565911 1501.644
```

```
In [4]: df = df.set_index(['districts','year'])
y = df['violent_cr']
X = df[['avg_temp','tot_rf']]
```

```
In [5]: #PooledOLS Estimation
X = sm.add_constant(X)
pols = PanelOLS(y,X)
pols_result = pols.fit()
print(pols_result.summary)
```

PanelOLS Estimation Summary

Dep. Variable:	violent_cr	R-squared:	0.0430
Estimator:	PanelOLS	R-squared (Between):	0.0967
No. Observations:	384	R-squared (Within):	-0.0184
Date:	Wed, Nov 12 2025	R-squared (Overall):	0.0430
Time:	18:28:44	Log-likelihood	-1289.0
Cov. Estimator:	Unadjusted	F-statistic:	8.5613
Entities:	32	P-value	0.0002
Avg Obs:	12.000	Distribution:	F(2,381)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	8.5613
		P-value	0.0002
Time periods:	12	Distribution:	F(2,381)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	7.1199	3.5388	2.0119	0.0449	0.1619	14.078
avg_temp	0.3714	0.1125	3.3021	0.0011	0.1503	0.5926
tot_rf	-0.0008	0.0008	-0.9385	0.3486	-0.0024	0.0009

```
In [6]: #FE Model Estimation
X = sm.add_constant(X)
FEmodel = PanelOLS(y,X,entity_effects=True)
feresult = FEmodel.fit()
print(feresult.summary)
```

PanelOLS Estimation Summary

Dep. Variable:	violent_cr	R-squared:	0.0027
Estimator:	PanelOLS	R-squared (Between):	-0.1497
No. Observations:	384	R-squared (Within):	0.0027
Date:	Wed, Nov 12 2025	R-squared (Overall):	-0.0786
Time:	18:28:44	Log-likelihood	-1150.5
Cov. Estimator:	Unadjusted	F-statistic:	0.4738
Entities:	32	P-value	0.6230
Avg Obs:	12.000	Distribution:	F(2, 350)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	0.4738
Time periods:	12	P-value	0.6230
Avg Obs:	32.000	Distribution:	F(2, 350)
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	23.220	12.266	1.8931	0.0592	-0.9031	47.344
avg_temp	-0.2709	0.4392	-0.6168	0.5378	-1.1348	0.5930
tot_rf	0.0003	0.0008	0.4033	0.6870	-0.0012	0.0018

F-test for Poolability: 11.930

P-value: 0.0000

Distribution: F(31,350)

Included effects: Entity

```
In [7]: #RE Model Estimation
from linearmodels.panel import RandomEffects
import statsmodels.api as sm
X = sm.add_constant(X)
REmodel = RandomEffects(y,X)
reresult = REmodel.fit()
print(reresult.summary)
```

RandomEffects Estimation Summary						
Dep. Variable:	violent_cr	R-squared:	0.0033			
Estimator:	RandomEffects	R-squared (Between):	0.0549			
No. Observations:	384	R-squared (Within):	-0.0016			
Date:	Wed, Nov 12 2025	R-squared (Overall):	0.0285			
Time:	18:28:44	Log-likelihood	-1167.7			
Cov. Estimator:	Unadjusted	F-statistic:	0.6212			
Entities:	32	P-value	0.5379			
Avg Obs:	12.000	Distribution:	F(2,381)			
Min Obs:	12.000	F-statistic (robust):	0.6212			
Max Obs:	12.000	P-value	0.5379			
Time periods:	12	Distribution:	F(2,381)			
Avg Obs:	32.000					
Min Obs:	32.000					
Max Obs:	32.000					
Parameter Estimates						
	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	8.9686	6.7143	1.3357	0.1824	-4.2332	22.170
avg_temp	0.2500	0.2346	1.0658	0.2872	-0.2112	0.7113
tot_rf	0.0005	0.0007	0.6711	0.5026	-0.0009	0.0019

In [8]: #Hausman Test

```

from numpy.linalg import inv
from scipy.stats import chi2

b_FE = feresult.params
b_RE = reresult.params

common_coef = list(set(b_FE.index) & set(b_RE.index))

if 'const' in common_coef:
    common_coef.remove('const')

b_FE = b_FE[common_coef]
b_RE = b_RE[common_coef]

V_FE = feresult.cov
V_RE = reresult.cov

diff = b_FE - b_RE
diff_var = V_FE.loc[common_coef, common_coef] - V_RE.loc[common_coef, common_coef]

hausman_stat = np.dot(np.dot(diff.T, inv(diff_var)), diff)

df_h = len(diff)
p_value = 1 - chi2.cdf(hausman_stat, df_h)

print("Hausman Test Statistic:", round(hausman_stat, 3))
print("Degrees of Freedom:", df_h)

```

```
print("p-value:", round(p_value, 4))
```

```
Hausman Test Statistic: 4.143  
Degrees of Freedom: 2  
p-value: 0.126
```

In [9]: *#Diagnostic Checks*

```
from statsmodels.stats.diagnostic import het_breushpagan, het_white  
from statsmodels.stats.stattools import durbin_watson
```

In [10]: *#Test for Heteroskedasticity*

```
#H0: No heteroskedasticity  
#H1: Heteroskedasticity exists
```

```
#p-value <= 0.05 ---> Heteroskedasticity; p-value > 0.05 ---> Homoskedasticity
```

```
print('Breusch-Pagan Test')  
residuals = reresult.resids  
bp_test = het_breushpagan(residuals, X)  
bp_labels = ['Lagrange multiplier statistic', 'p-value', 'f-value', 'f p-value']  
print(dict(zip(bp_labels, bp_test)))  
print()  
print('White Test')  
white_test = het_white(residuals, X)  
white_labels = ['LM stat', 'LM p-value', 'F p-value']  
print(dict(zip(white_labels, white_test)))
```

Breusch-Pagan Test
{'Lagrange multiplier statistic': np.float64(4.210322820038144), 'p-value': np.float64(0.12182600874362186), 'f-value': np.float64(2.1118701887129188), 'f p-value': np.float64(0.12242582224810573)}

White Test

```
{'LM stat': np.float64(13.773621666102258), 'LM p-value': np.float64(0.017113181383934493), 'F p-value': np.float64(2.8125651193287555)}
```

In [11]: *#Test for serial correlation (autocorrelation)*

```
#Durbin-Watson statistic ranges between 0 to 4
```

```
#DW statistic = 2 ---> No autocorrelation  
#DW statistic < 2 ---> Positive autocorrelation  
#DW statistic > 2 ---> Negative autocorrelation
```

```
print('Durbin-Watson Test')  
dw_value = durbin_watson(residuals)  
print("Durbin-Watson statistic: ", round(dw_value, 3))
```

Durbin-Watson Test

Durbin-Watson statistic: 1.011

In [12]: *from scipy import stats*

```
#Test for cross-section dependency
```

```
#H0: No cross-section dependency
```

```
#H1: Cross-section dependency exists

print('Breusch-Pagan LM Test')
resid_df = residuals.unstack(level=0)
T = resid_df.shape[0]
N = resid_df.shape[1]

rho = resid_df.corr().values
upper_tri_idx = np.triu_indices(N, k=1)
rho_upper = rho[upper_tri_idx]
LM_stat = T * np.sum(rho_upper**2)
p_value = 1 - stats.chi2.cdf(LM_stat, N*(N-1)/2)

print(f"Breusch-Pagan LM statistic: {LM_stat:.3f}")
print(f"p-value: {p_value:.4f}")
print()

print('Pesaran CD Test')
CD_stat = np.sqrt(2 / (N*(N-1))) * np.sum(rho_upper)
p_value_cd = 2 * (1 - stats.norm.cdf(abs(CD_stat)))

print(f"Pesaran CD statistic: {CD_stat:.3f}")
print(f"p-value: {p_value_cd:.4f}")
```

```
Breusch-Pagan LM Test
Breusch-Pagan LM statistic: 1638.685
p-value: 0.0000
```

```
Pesaran CD Test
Pesaran CD statistic: 8.856
p-value: 0.0000
```

```
In [13]: #Re-estimate RE Model
```

```
In [14]: #RE with cov.type 'clustered'
re_model_robust1 = REmodel.fit(cov_type='clustered', cluster_entity=True)
print(re_model_robust1.summary)
```

RandomEffects Estimation Summary

Dep. Variable:	violent_cr	R-squared:	0.0033
Estimator:	RandomEffects	R-squared (Between):	0.0549
No. Observations:	384	R-squared (Within):	-0.0016
Date:	Wed, Nov 12 2025	R-squared (Overall):	0.0285
Time:	18:28:44	Log-likelihood	-1167.7
Cov. Estimator:	Clustered	F-statistic:	0.6212
Entities:	32	P-value	0.5379
Avg Obs:	12.000	Distribution:	F(2, 381)
Min Obs:	12.000		
Max Obs:	12.000	F-statistic (robust):	0.7032
		P-value	0.4956
Time periods:	12	Distribution:	F(2, 381)
Avg Obs:	32.000		
Min Obs:	32.000		
Max Obs:	32.000		

Parameter Estimates

	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	8.9686	6.1474	1.4589	0.1454	-3.1186	21.056
avg_temp	0.2500	0.2123	1.1777	0.2397	-0.1674	0.6674
tot_rf	0.0005	0.0010	0.4719	0.6373	-0.0015	0.0025

```
In [15]: #RE with cov.type 'kernel' (Driscoll-Kraay Method)
re_model_robust2 = REmodel.fit(cov_type='kernel')
print(re_model_robust2.summary)
```

RandomEffects Estimation Summary						
Dep. Variable:	violent_cr	R-squared:	0.0033			
Estimator:	RandomEffects	R-squared (Between):	0.0549			
No. Observations:	384	R-squared (Within):	-0.0016			
Date:	Wed, Nov 12 2025	R-squared (Overall):	0.0285			
Time:	18:28:44	Log-likelihood	-1167.7			
Cov. Estimator:	Driscoll-Kraay					
		F-statistic:	0.6212			
Entities:	32	P-value	0.5379			
Avg Obs:	12.000	Distribution:	$F(2, 381)$			
Min Obs:	12.000					
Max Obs:	12.000	F-statistic (robust):	1.0834			
		P-value	0.3395			
Time periods:	12	Distribution:	$F(2, 381)$			
Avg Obs:	32.000					
Min Obs:	32.000					
Max Obs:	32.000					
Parameter Estimates						
	Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI
const	8.9686	8.3392	1.0755	0.2828	-7.4280	25.365
avg_temp	0.2500	0.2355	1.0615	0.2891	-0.2131	0.7131
tot_rf	0.0005	0.0010	0.4633	0.6434	-0.0016	0.0025

```
In [16]: # Check residuals and fitted values
df['residuals1'] = re_model_robust1.resids
df['fitted1'] = re_model_robust1.fitted_values

import matplotlib.pyplot as plt

plt.scatter(df['fitted1'], df['residuals1'], alpha=0.6)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values (RE model with Clustered Standard Errors)')
plt.show()

sm.qqplot(df['residuals1'], line='45', fit=True)
plt.title('Q-Q Plot of Residuals')
plt.show()

plt.hist(df['residuals1'], bins=30, edgecolor='black', alpha=0.7)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (RE model)')
plt.show()

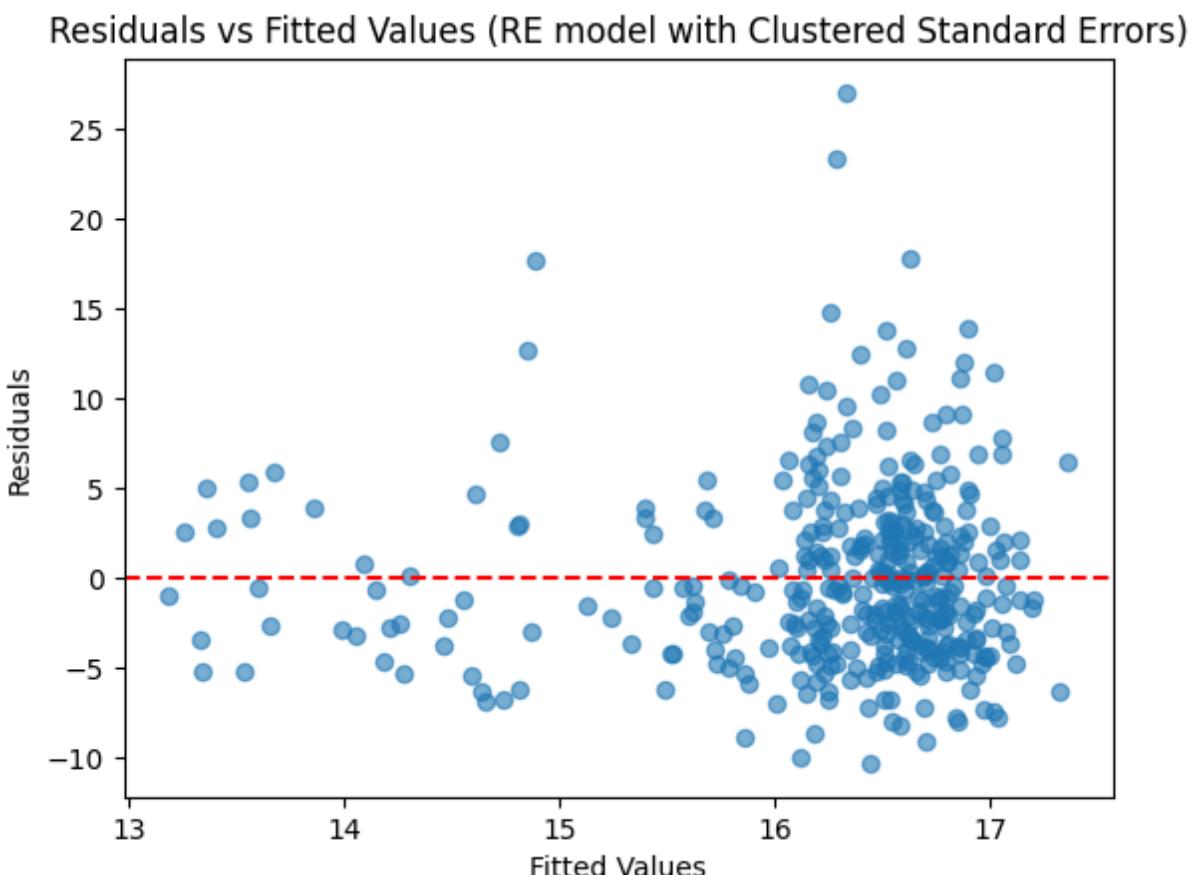
resid_df = df['residuals1'].unstack(level=0)
plt.plot(resid_df.mean(axis=1))
plt.title('Average Residuals over Time')
plt.xlabel('Year')
plt.ylabel('Mean Residual')
```

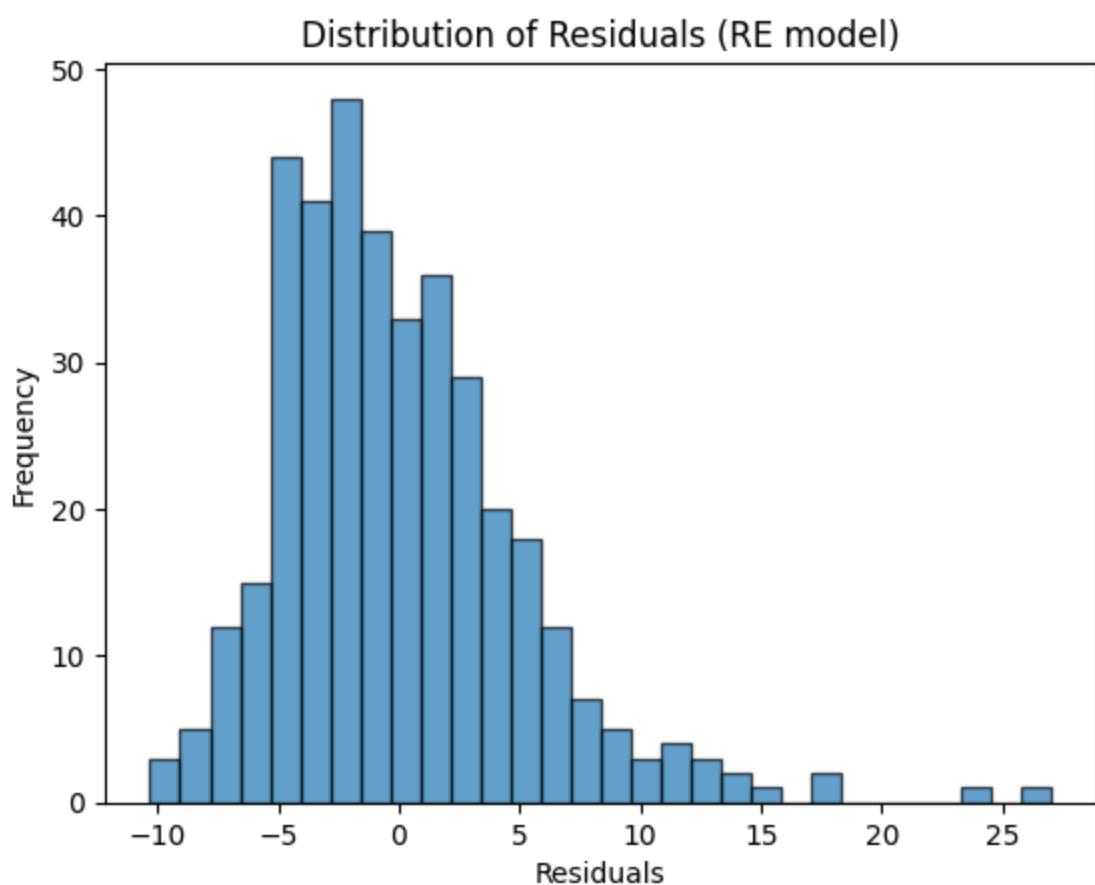
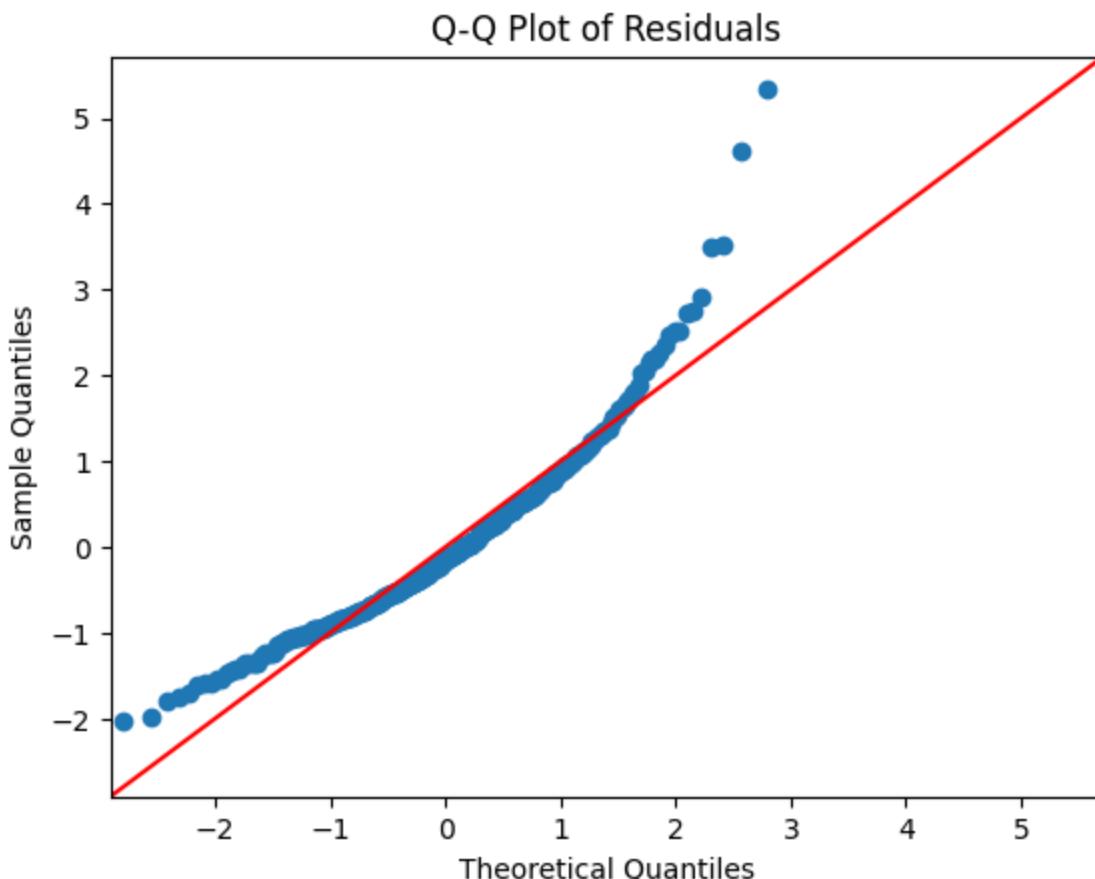
```
plt.show()

from scipy.stats import shapiro

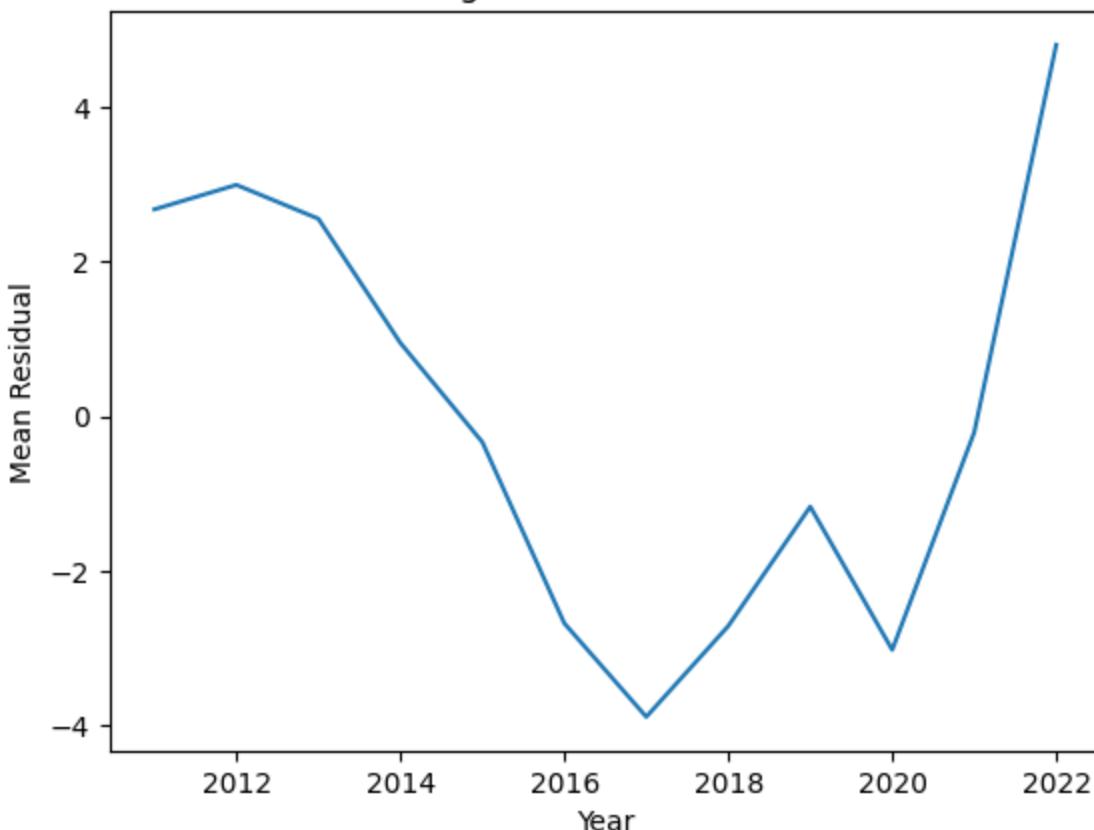
#Test for normality

stat, p = shapiro(df['residuals1'])
print(f"Shapiro-Wilk Test: Statistic={stat:.3f}, p-value={p:.4f}")
```





Average Residuals over Time



Shapiro-Wilk Test: Statistic=0.932, p-value=0.0000

```
In [17]: # Check residuals and fitted values
df['residuals2'] = re_model_robust2.resids
df['fitted2'] = re_model_robust2.fitted_values

import matplotlib.pyplot as plt

plt.scatter(df['fitted2'], df['residuals2'], alpha=0.6)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values (RE model with Driscoll-Kraay)')
plt.show()

sm.qqplot(df['residuals2'], line='45', fit=True)
plt.title('Q-Q Plot of Residuals')
plt.show()

plt.hist(df['residuals2'], bins=30, edgecolor='black', alpha=0.7)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (RE model)')
plt.show()

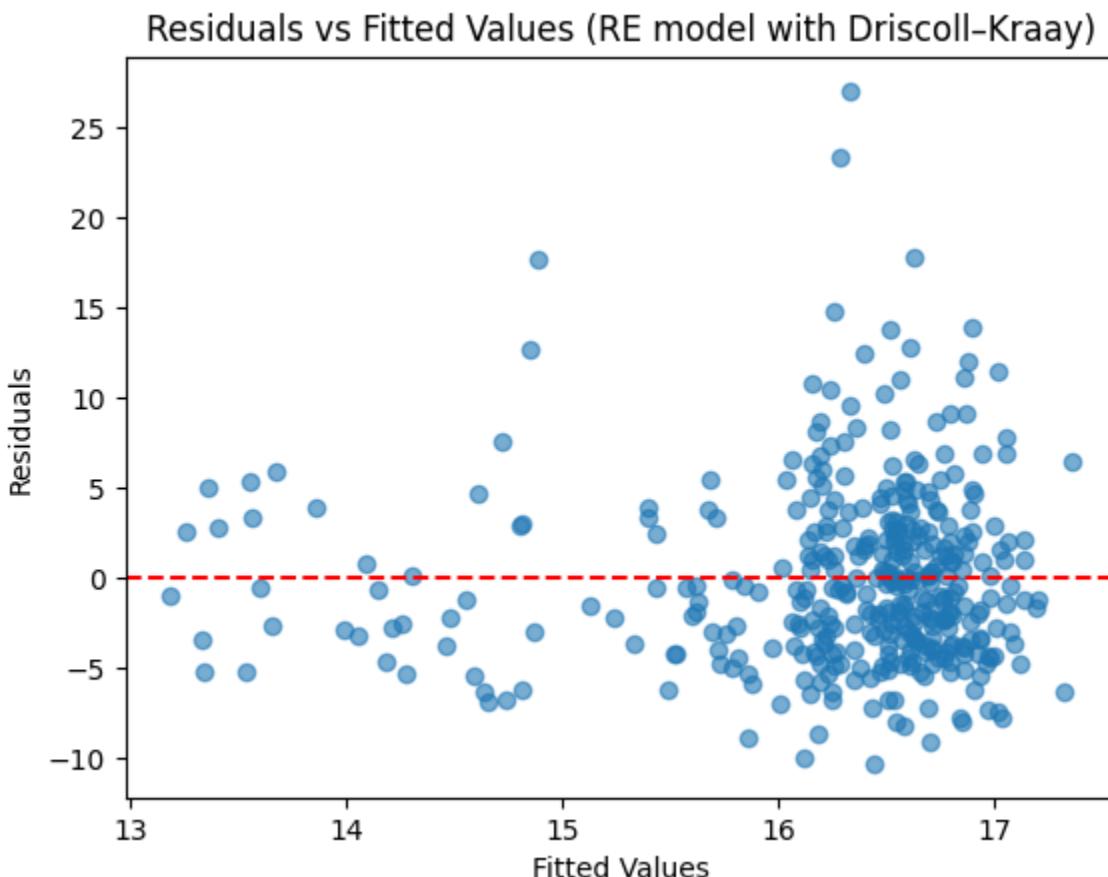
resid_df = df['residuals2'].unstack(level=0)
plt.plot(resid_df.mean(axis=1))
plt.title('Average Residuals over Time')
plt.xlabel('Year')
```

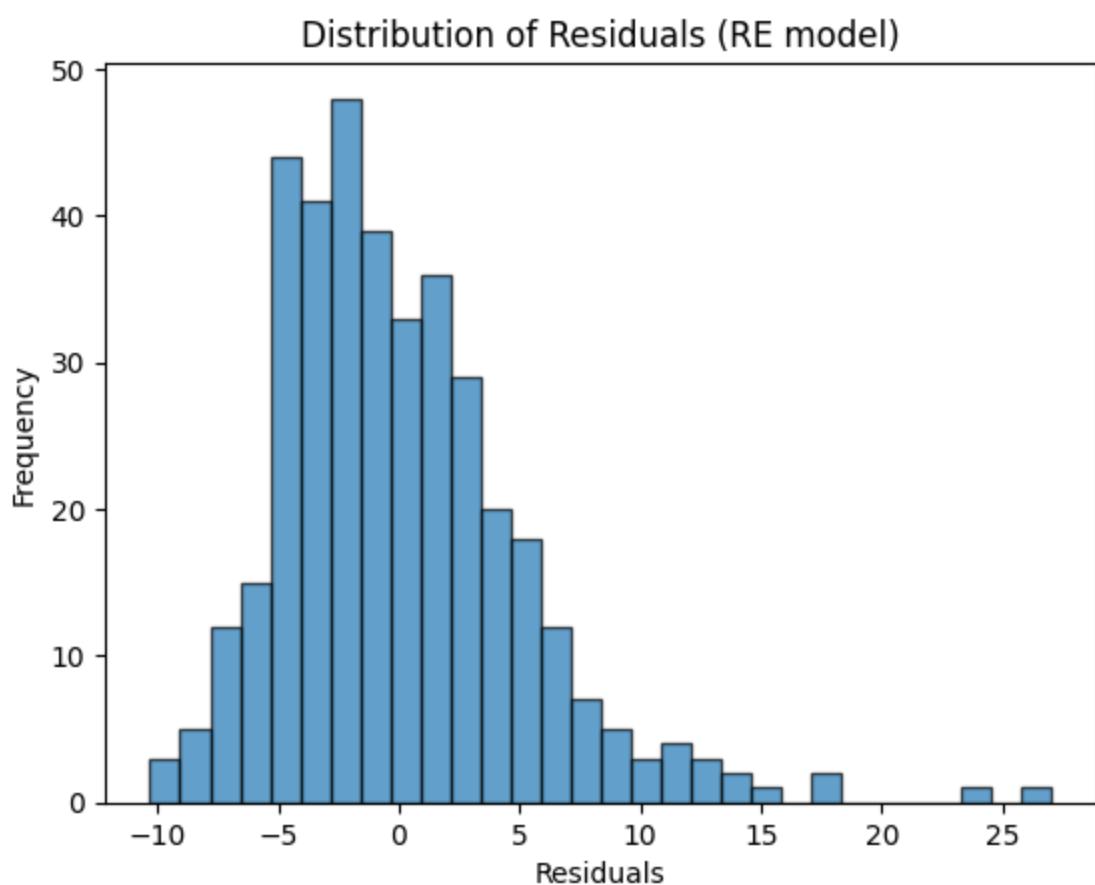
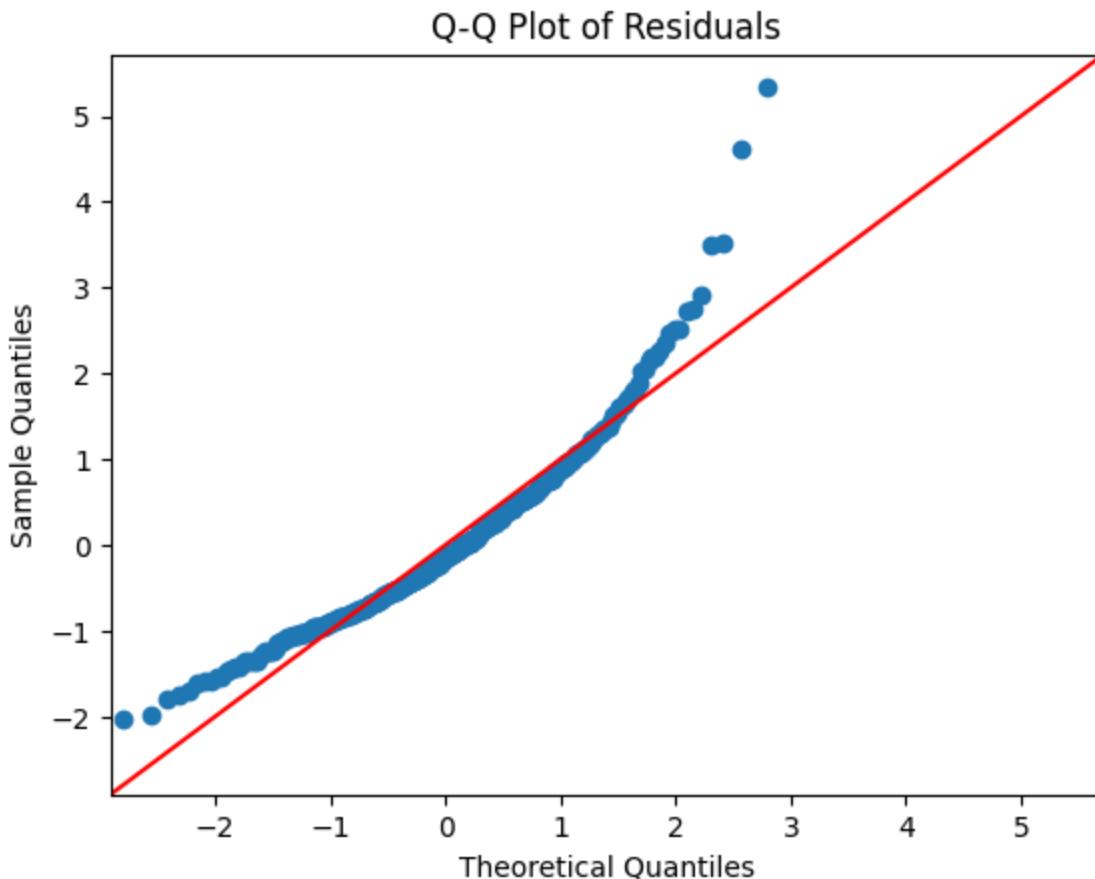
```
plt.ylabel('Mean Residual')
plt.show()

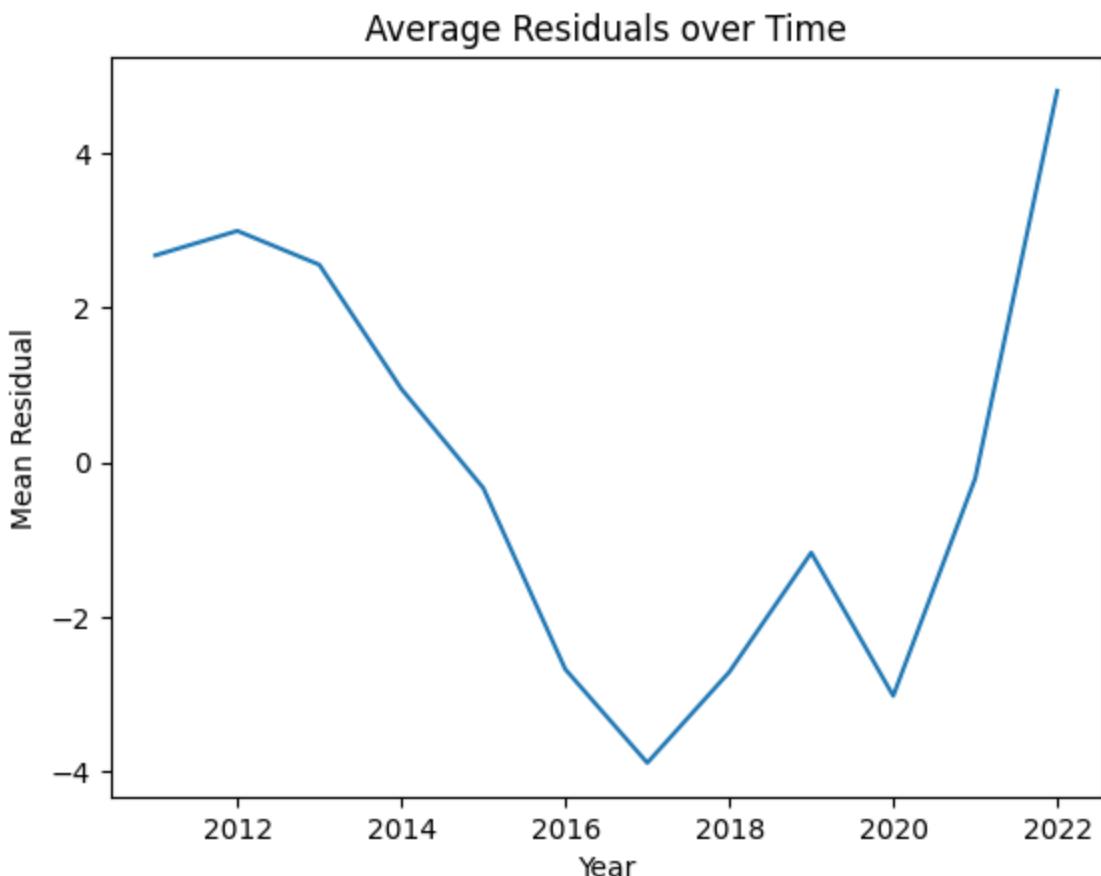
from scipy.stats import shapiro

#Test for normality

stat, p = shapiro(df['residuals2'])
print(f"Shapiro-Wilk Test: Statistic={stat:.3f}, p-value={p:.4f}")
```







Shapiro-Wilk Test: Statistic=0.932, p-value=0.0000

In []: